

```

import pandas as pd

train_df = pd.read_csv('/content/customer_churn_dataset-training-master[1].csv')
test_df = pd.read_csv('/content/customer_churn_dataset-testing-master[1].csv')

print("Training DataFrame - Head:")
print(train_df.head())

```

Training DataFrame - Head:

	CustomerID	Age	Gender	Tenure	Usage Frequency	Support Calls	\
0	2.0	30.0	Female	39.0	14.0	5.0	
1	3.0	65.0	Female	49.0	1.0	10.0	
2	4.0	55.0	Female	14.0	4.0	6.0	
3	5.0	58.0	Male	38.0	21.0	7.0	
4	6.0	23.0	Male	32.0	20.0	5.0	

	Payment Delay	Subscription Type	Contract Length	Total Spend	\
0	18.0	Standard	Annual	932.0	
1	8.0	Basic	Monthly	557.0	
2	18.0	Basic	Quarterly	185.0	
3	7.0	Standard	Monthly	396.0	
4	8.0	Basic	Monthly	617.0	

	Last Interaction	Churn
0	17.0	1.0
1	6.0	1.0
2	3.0	1.0
3	29.0	1.0
4	20.0	1.0

```

print("\nTraining DataFrame - Info:")
train_df.info()

```

```

print("\nTraining DataFrame - Description:")
print(train_df.describe())

```

```

print("\nTraining DataFrame - Missing Values:")
print(train_df.isnull().sum())

```

2	Gender	440832 non-null	object
3	Tenure	440832 non-null	float64
4	Usage Frequency	440832 non-null	float64
5	Support Calls	440832 non-null	float64
6	Payment Delay	440832 non-null	float64
7	Subscription Type	440832 non-null	object
8	Contract Length	440832 non-null	object
9	Total Spend	440832 non-null	float64
10	Last Interaction	440832 non-null	float64
11	Churn	440832 non-null	float64

dtypes: float64(9), object(3)

memory usage: 40.4+ MB

Training DataFrame - Description:

	CustomerID	Age	Tenure	Usage Frequency	\
count	440832.000000	440832.000000	440832.000000	440832.000000	

mean 225398.667955 39.373153 31.256336 15.807494

std 129531.918550 12.442369 17.255727 8.586242

min 2.000000 18.000000 1.000000 1.000000

25% 113621.750000 29.000000 16.000000 9.000000

50% 226125.500000 39.000000 32.000000 16.000000

75% 337739.250000 48.000000 46.000000 23.000000

max 449999.000000 65.000000 60.000000 30.000000

	Support Calls	Payment Delay	Total Spend	Last Interaction	\
--	---------------	---------------	-------------	------------------	---

count 440832.000000 440832.000000 440832.000000 440832.000000

mean 3.604437 12.965722 631.616223 14.480868

std 3.070218 8.258063 240.803001 8.596208

```
min          0.000000
25%         0.000000
50%         1.000000
75%         1.000000
max          1.000000
```

```
Training DataFrame - Missing Values:
CustomerID      1
Age             1
Gender           1
Tenure           1
Usage Frequency  1
Support Calls    1
Payment Delay    1
Subscription Type 1
Contract Length   1
Total Spend      1
Last Interaction  1
Churn            1
dtype: int64
```

```
print("\nTesting DataFrame - Head:")
print(test_df.head())

print("\nTesting DataFrame - Info:")
test_df.info()

print("\nTesting DataFrame - Description:")
print(test_df.describe())

print("\nTesting DataFrame - Missing Values:")
print(test_df.isnull().sum())
```

```
2   Gender          64374 non-null  object
3   Tenure          64374 non-null  int64
4   Usage Frequency 64374 non-null  int64
5   Support Calls   64374 non-null  int64
6   Payment Delay   64374 non-null  int64
7   Subscription Type 64374 non-null  object
8   Contract Length 64374 non-null  object
9   Total Spend     64374 non-null  int64
10  Last Interaction 64374 non-null  int64
11  Churn           64374 non-null  int64
dtypes: int64(9), object(3)
memory usage: 5.9+ MB
```

Testing DataFrame - Description:

```
CustomerID      0  
Age            0  
Gender         0  
Tenure          0  
Usage Frequency 0  
Support Calls   0  
Payment Delay    0  
Subscription Type 0  
Contract Length   0  
Total Spend      0  
Last Interaction  0  
Churn           0  
dtype: int64
```

```
print("Shape of train_df before dropping missing values:", train_df.shape)  
train_df.dropna(inplace=True)  
print("Shape of train_df after dropping missing values:", train_df.shape)  
  
print("\nMissing values in train_df after dropping:")  
print(train_df.isnull().sum())
```

```
Shape of train_df before dropping missing values: (440833, 12)  
Shape of train_df after dropping missing values: (440832, 12)
```

```
Missing values in train_df after dropping:
```

```
CustomerID      0  
Age            0  
Gender         0  
Tenure          0  
Usage Frequency 0  
Support Calls   0  
Payment Delay    0  
Subscription Type 0  
Contract Length   0  
Total Spend      0  
Last Interaction  0  
Churn           0  
dtype: int64
```

```
import pandas as pd  
  
categorical_cols = ['Gender', 'Subscription Type', 'Contract Length']  
  
# Apply One-Hot Encoding to training data  
train_df_encoded = pd.get_dummies(train_df, columns=categorical_cols, drop_first=False)  
print("Shape of train_df after One-Hot Encoding:", train_df_encoded.shape)  
  
# Apply One-Hot Encoding to testing data, ensuring consistent columns with training data  
test_df_encoded = pd.get_dummies(test_df, columns=categorical_cols, drop_first=False)  
  
# Align columns - this is crucial to ensure both dataframes have the same columns after encoding  
train_cols = set(train_df_encoded.columns)  
test_cols = set(test_df_encoded.columns)  
  
missing_in_test = list(train_cols - test_cols)  
for col in missing_in_test:  
    test_df_encoded[col] = 0  
  
missing_in_train = list(test_cols - train_cols)  
for col in missing_in_train:  
    train_df_encoded[col] = 0  
  
test_df_encoded = test_df_encoded[train_df_encoded.columns]  
  
train_df = train_df_encoded  
test_df = test_df_encoded  
  
print("Shape of test_df after One-Hot Encoding and column alignment:", test_df.shape)  
print("First 5 rows of train_df after encoding:")  
print(train_df.head())
```

```
Shape of train_df after One-Hot Encoding: (440832, 17)  
Shape of test_df after One-Hot Encoding and column alignment: (64374, 17)  
First 5 rows of train_df after encoding:  
CustomerID  Age  Tenure  Usage Frequency  Support Calls  Payment Delay  \  
0          2.0  30.0     39.0             14.0           5.0          18.0
```

1	3.0	65.0	49.0	1.0	10.0	8.0
2	4.0	55.0	14.0	4.0	6.0	18.0
3	5.0	58.0	38.0	21.0	7.0	7.0
4	6.0	23.0	32.0	20.0	5.0	8.0

	Total Spend	Last Interaction	Churn	Gender_Female	Gender_Male	\
0	932.0	17.0	1.0	True	False	
1	557.0	6.0	1.0	True	False	
2	185.0	3.0	1.0	True	False	
3	396.0	29.0	1.0	False	True	
4	617.0	20.0	1.0	False	True	

	Subscription Type_Basic	Subscription Type_Premium	\
0	False	False	
1	True	False	
2	True	False	
3	False	False	
4	True	False	

	Subscription Type_Standard	Contract Length_Annual	\
0	True	True	
1	False	False	
2	False	False	
3	True	False	
4	False	False	

	Contract Length_Monthly	Contract Length_Quarterly
0	False	False
1	True	False
2	False	True
3	True	False
4	True	False

```

from sklearn.preprocessing import StandardScaler

# Identify numerical columns to be scaled (excluding CustomerID and Churn for now)
numerical_cols = ['Age', 'Tenure', 'Usage Frequency', 'Support Calls', 'Payment Delay', 'Total Spend', 'Last In-
 
# Initialize StandardScaler
scaler = StandardScaler()

# Apply StandardScaler to training data (fit and transform)
train_df[numerical_cols] = scaler.fit_transform(train_df[numerical_cols])

# Apply StandardScaler to testing data (transform only)
test_df[numerical_cols] = scaler.transform(test_df[numerical_cols])

# Drop 'CustomerID' column from both dataframes
train_df.drop('CustomerID', axis=1, inplace=True)
test_df.drop('CustomerID', axis=1, inplace=True)

print("First 5 rows of train_df after scaling numerical features and dropping CustomerID:")
print(train_df.head())
print("\nFirst 5 rows of test_df after scaling numerical features and dropping CustomerID:")
print(test_df.head())

```

4	-0.060698	0.642043	1.0	False	True
---	-----------	----------	-----	-------	------

	Subscription Type_Basic	Subscription Type_Premium	\
0	False	False	
1	True	False	
2	True	False	
3	False	False	
4	True	False	

```

First 5 rows of test_df after scaling numerical features and dropping CustomerID:
   Age   Tenure  Usage_Frequency  Support_Calls  Payment_Delay \
0 -1.396291 -0.362566      -0.210511      0.128839      1.699466
1  0.130751 -0.188711      1.420007      1.105969      0.004151
2  0.612974 -0.246662     -0.676373     -0.522581      1.941654
3 -0.351473 -1.289796     -0.443442      0.454549      0.488527
4  1.095198  1.549845      0.954145      1.757390     -1.327882

   Total_Spend  Last_Interaction  Churn  Gender_Female  Gender_Male \
0      -0.139601      -0.637592       1      True      False
1      -0.197740       0.642043       0      True      False
2       0.520691       0.758374       0     False      True
3      -1.659517       0.409382       0     False      True
4      -0.409531       0.409382       0      True      False

   Subscription_Type_Basic  Subscription_Type_Premium \
0                  True                  False
1                 False                  False
2                 False                  True
3                 False                  True
4                 False                  False

   Subscription_Type_Standard  Contract_Length_Annual \
0                  False                  False
1                  True                  False
2                 False                  True
3                 False                  False
4                  True                  True

   Contract_Length_Monthly  Contract_Length_Quarterly
0                  True                  False
1                  True                  False
2                 False                  False
3                 False                  True
4                 False                  False

```

```

X = train_df.drop('Churn', axis=1)
y = train_df['Churn']

print("Shape of X (features):")
print(X.shape)
print("Shape of y (target):")
print(y.shape)

```

```

Shape of X (features):
(440832, 15)
Shape of y (target):
(440832,)

```

```

from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

print("Shape of X_train:", X_train.shape)
print("Shape of X_val:", X_val.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_val:", y_val.shape)

Shape of X_train: (352665, 15)
Shape of X_val: (88167, 15)
Shape of y_train: (352665,)
Shape of y_val: (88167,)

```

```

import matplotlib.pyplot as plt
import seaborn as sns

print("Libraries matplotlib.pyplot and seaborn imported successfully.")

Libraries matplotlib.pyplot and seaborn imported successfully.

```

```

plt.figure(figsize=(16, 12))
sns.heatmap(train_df.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of All Features in Training Data')
plt.show()

```

Correlation Matrix of All Features in Training

	Age	Tenure	Usage Frequency	Support Calls	Payment Delay	Total Spend	Last Interaction	Churn	Gender_Female	Gender_Male	Subscription Type_Basic
Age	1.00	-0.01	-0.01	0.16	0.06	-0.08	0.03	0.22	0.03	-0.03	0.01
Tenure	-0.01	1.00	-0.03	-0.03	-0.02	0.02	-0.01	-0.05	-0.01	0.01	0.01
Usage Frequency	-0.01	-0.03	1.00	-0.02	-0.01	0.02	-0.00	-0.05	-0.01	0.01	0.01
Support Calls	0.16	-0.03	-0.02	1.00	0.16	-0.22	0.08	0.57	0.09	-0.09	0.01
Payment Delay	0.06	-0.02	-0.01	0.16	1.00	-0.12	0.04	0.31	0.05	-0.05	0.01
Total Spend	-0.08	0.02	0.02	-0.22	-0.12	1.00	-0.06	-0.43	-0.07	0.07	-0.01
Last Interaction	0.03	-0.01	-0.00	0.08	0.04	-0.06	1.00	0.15	-0.13	0.13	0.01
Churn	0.22	-0.05	-0.05	0.57	0.31	-0.43	0.15	1.00	0.18	-0.18	0.01
Gender_Female	0.03	-0.01	-0.01	0.09	0.05	-0.07	-0.13	0.18	1.00	-1.00	0.01
Gender_Male	-0.03	0.01	0.01	-0.09	-0.05	0.07	0.13	-0.18	-1.00	1.00	-0.01
Subscription Type_Basic	0.00	0.03	0.00	0.01	0.01	-0.01	0.00	0.02	0.00	-0.00	1.00
Subscription Type_Premium	-0.00	-0.02	0.00	-0.01	-0.00	0.00	-0.00	-0.01	-0.00	0.00	-0.01
Subscription Type_Standard	-0.00	-0.01	-0.00	-0.00	-0.00	0.00	-0.00	-0.01	0.00	-0.00	-0.01
Contract Length_Annual	-0.03	0.01	0.01	-0.09	-0.05	0.07	-0.02	-0.18	-0.03	0.03	-0.01
Contract Length_Monthly	0.09	-0.02	-0.02	0.22	0.12	-0.17	0.06	0.43	0.07	-0.07	0.01
Contract Length_Quarterly	-0.04	0.01	0.01	-0.09	-0.05	0.07	-0.02	-0.18	-0.03	0.03	-0.01
Age											
Tenure											
Usage Frequency											
Support Calls											
Payment Delay											
Total Spend											
Last Interaction											
Churn											
Gender_Female											
Gender_Male											
Subscription Type_Basic											

Interpretation of Correlation Heatmap:

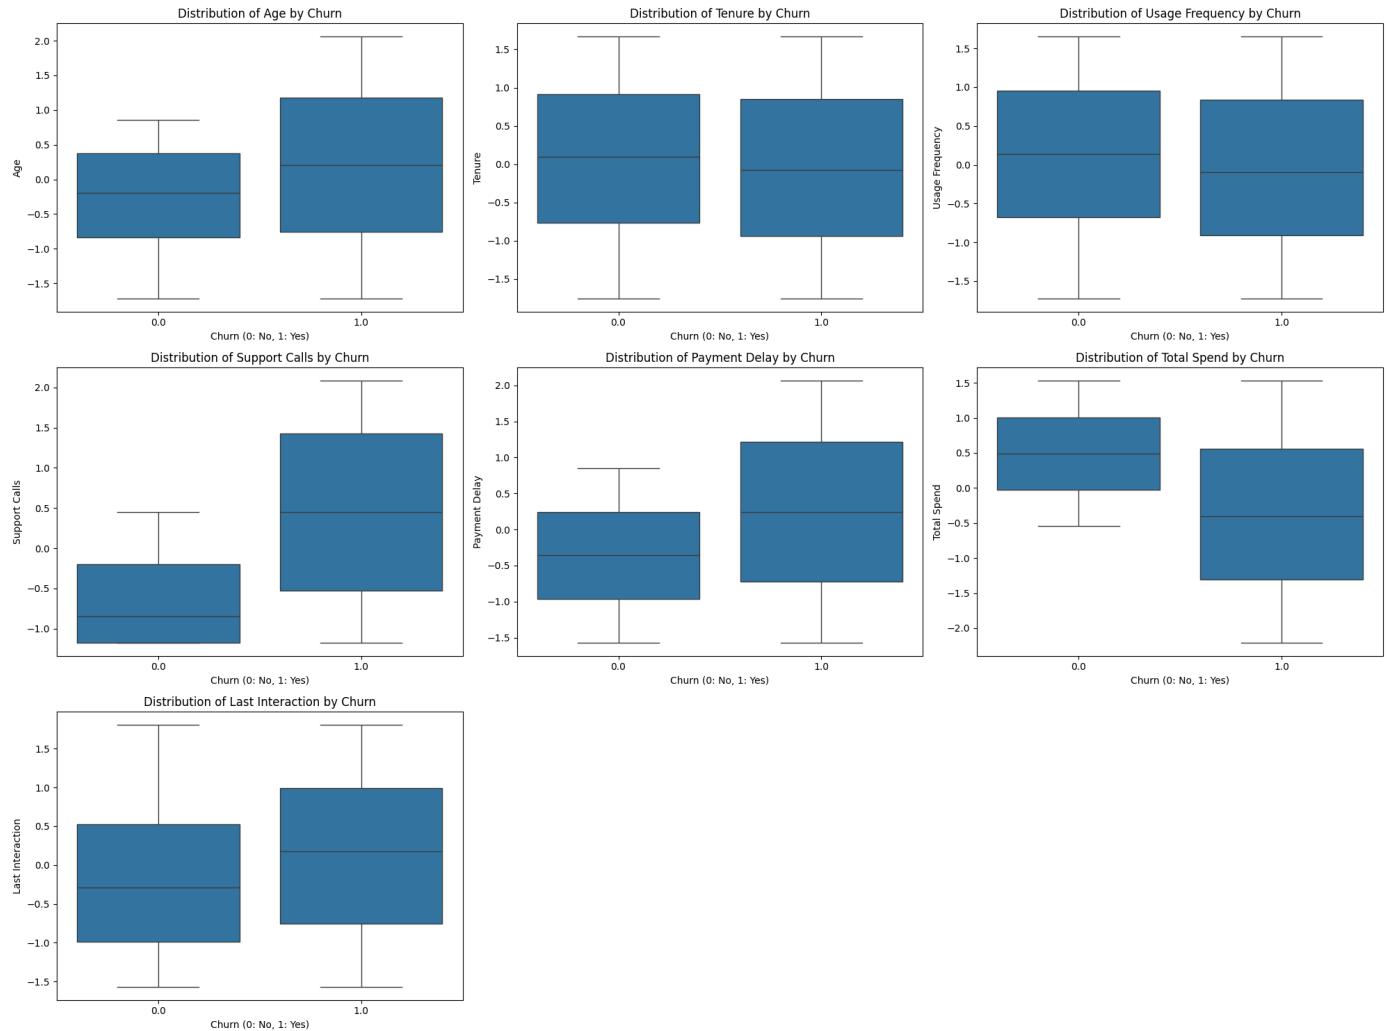
The correlation heatmap provides insights into the linear relationships between all features, including the 'Churn' feature.

- Positive Correlation with Churn:** Features like `Payment Delay`, `Support Calls`, and `Usage Frequency` show a positive correlation with 'Churn'.
- Negative Correlation with Churn:** `Tenure` and `Total Spend` exhibit moderate negative correlations with 'Churn'.
- Inter-feature Correlations:** There are some notable correlations between independent features. For instance, `Age` and `Tenure` have a very slight negative correlation.
- Categorical Features:** The one-hot encoded categorical features (Gender, Subscription Type, Contract Length) show no significant correlations with other features.

This heatmap serves as a good starting point for further investigation into specific relationships with box plots a

```
numerical_cols_for_boxplot = ['Age', 'Tenure', 'Usage Frequency', 'Support Calls', 'Payment Delay', 'Total Spend']

plt.figure(figsize=(20, 15))
for i, col in enumerate(numerical_cols_for_boxplot):
    plt.subplot(3, 3, i + 1) # Adjust subplot grid based on number of numerical columns
    sns.boxplot(x='Churn', y=col, data=train_df)
    plt.title(f'Distribution of {col} by Churn')
    plt.xlabel('Churn (0: No, 1: Yes)')
    plt.ylabel(col)
plt.tight_layout()
plt.show()
```



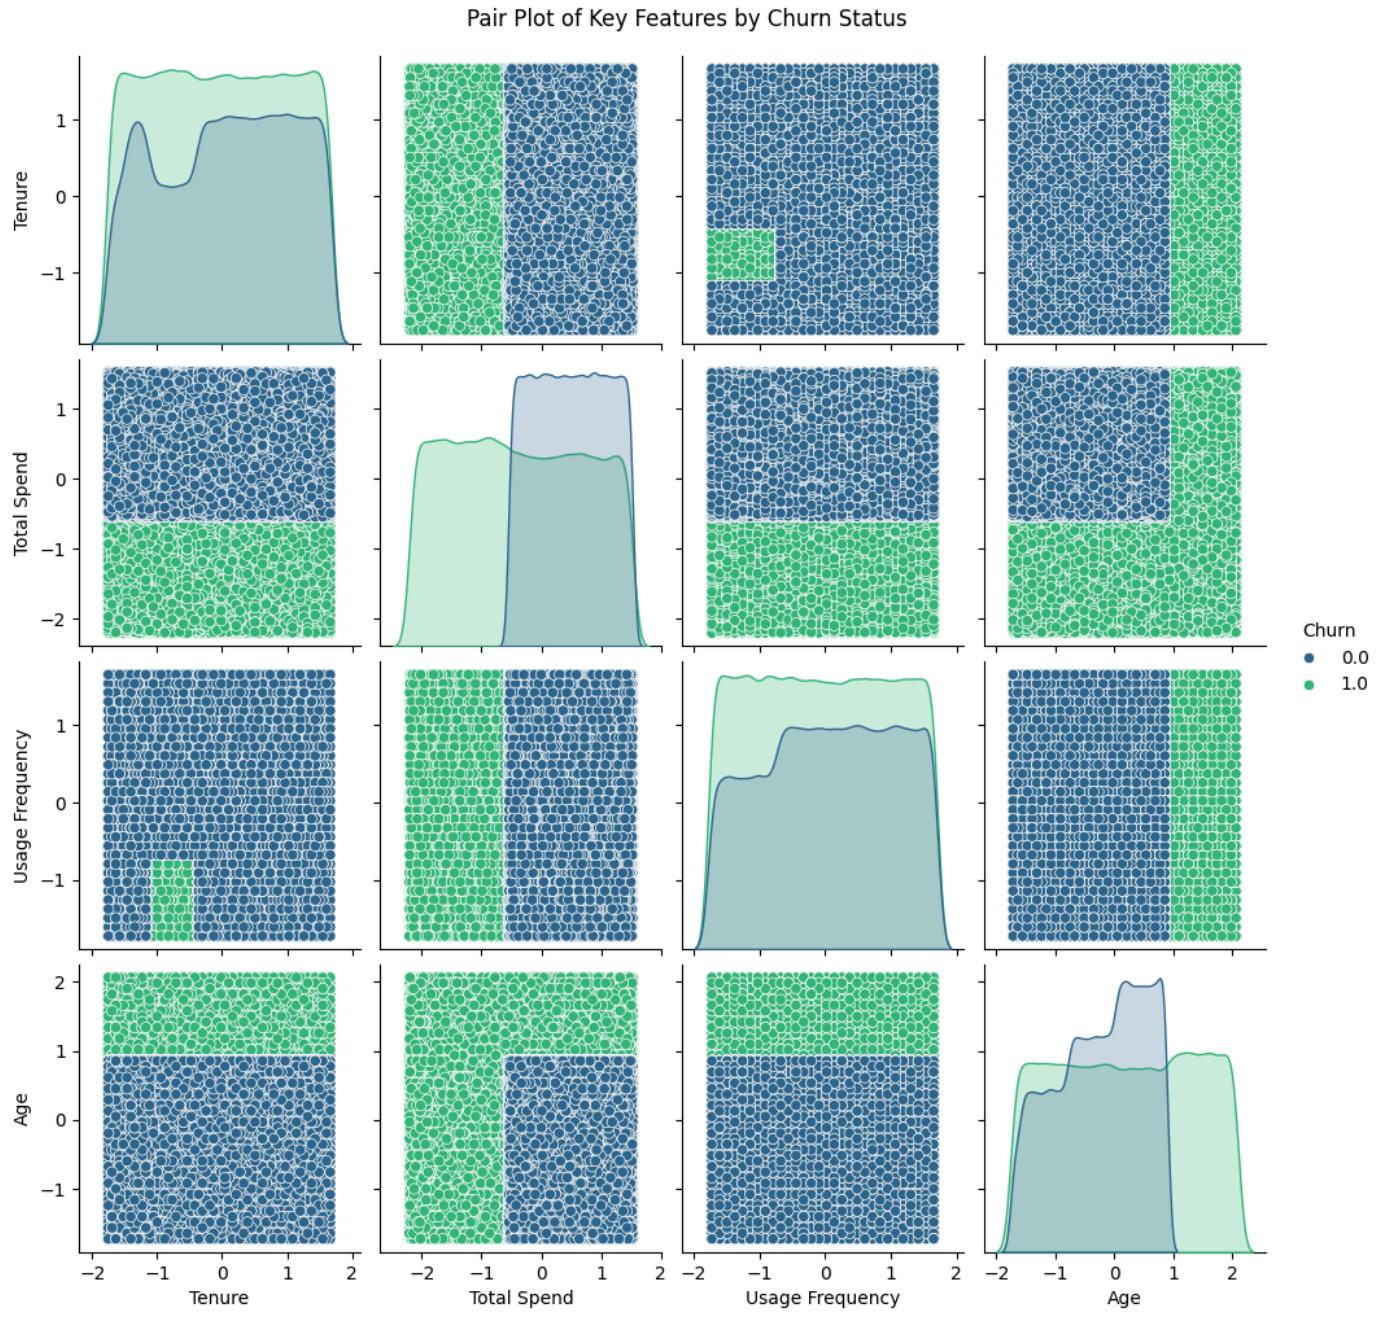
Interpretation of Box Plots (Numerical Features vs. Churn):

The box plots illustrate the distribution of each numerical feature for customers who churned (1) and those who did

- ****Age**:** The median age appears similar for both churned and non-churned customers, indicating age might not be a significant factor.
- ****Tenure**:** Non-churned customers tend to have a higher median tenure and a wider distribution towards longer tenures compared to churned customers.
- ****Usage Frequency**:** Churned customers seem to have a slightly higher median and range for usage frequency, which may indicate more frequent usage.
- ****Support Calls**:** Churned customers exhibit a higher median and range for support calls, implying that frequent support interactions are associated with churn.
- ****Payment Delay**:** There's a noticeable difference, with churned customers generally having higher payment delays.
- ****Total Spend**:** Non-churned customers show a higher median total spend and a larger interquartile range, indicating higher overall spending.
- ****Last Interaction**:** The distributions for last interaction seem somewhat similar between churned and non-churned customers.

Overall, features like `Tenure`, `Payment Delay`, `Support Calls`, and `Total Spend` show distinct differences in their distributions between churned and non-churned customers.

```
key_features_for_pairplot = ['Tenure', 'Total Spend', 'Usage Frequency', 'Age', 'Churn']
sns.pairplot(train_df[key_features_for_pairplot], hue='Churn', palette='viridis', diag_kind='kde')
plt.suptitle('Pair Plot of Key Features by Churn Status', y=1.02) # Adjust suptitle position
plt.show()
```



Interpretation of Pair Plots:

The pair plots provide a matrix of scatter plots for each combination of the selected key numerical features, with

- **Tenure vs. Total Spend**: There's a clear positive correlation, as expected. Customers with higher `Tenure` generally have higher `Total Spend`.
- **Usage Frequency vs. Tenure/Total Spend**: `Usage Frequency` doesn't show a very strong linear relationship with either `Tenure` or `Total Spend`.
- **Age vs. Other Features**: `Age` generally shows a mild positive correlation with `Tenure` and `Total Spend`. The relationship with `Usage Frequency` is less clear.
- **Diagonal (KDE Plots)**: The diagonal plots show the density distribution for each feature, separated by churn status.

Overall, the pair plots confirm that `Tenure` and `Total Spend` are important indicators, with lower values in both categories associated with higher churn rates.

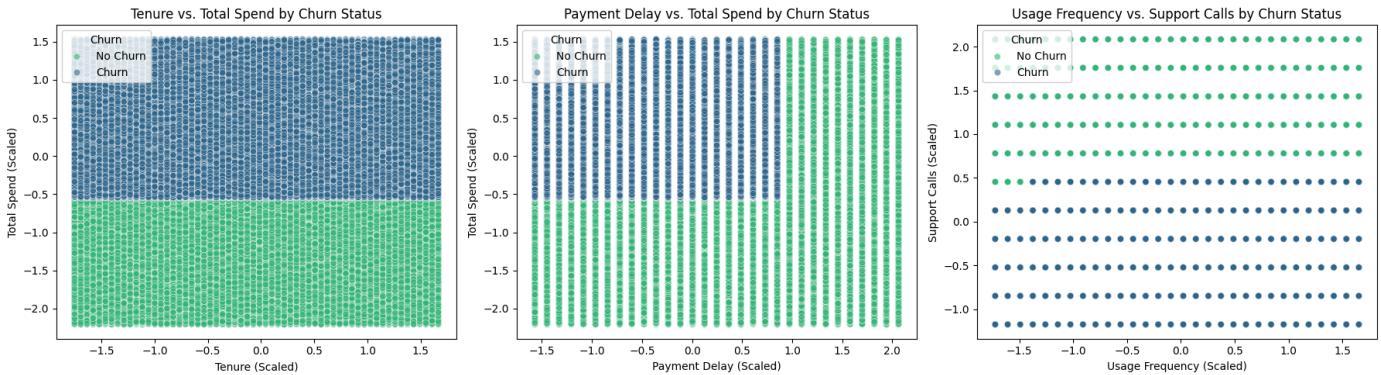
```
plt.figure(figsize=(18, 5))

# Scatter Plot 1: Tenure vs. Total Spend
plt.subplot(1, 3, 1)
sns.scatterplot(x='Tenure', y='Total Spend', hue='Churn', data=train_df, palette='viridis', alpha=0.6)
plt.title('Tenure vs. Total Spend by Churn Status')
plt.xlabel('Tenure (Scaled)')
plt.ylabel('Total Spend (Scaled)')
plt.legend(title='Churn', loc='upper left', labels=['No Churn', 'Churn'])

# Scatter Plot 2: Payment Delay vs. Total Spend
plt.subplot(1, 3, 2)
sns.scatterplot(x='Payment Delay', y='Total Spend', hue='Churn', data=train_df, palette='viridis', alpha=0.6)
plt.title('Payment Delay vs. Total Spend by Churn Status')
plt.xlabel('Payment Delay (Scaled)')
plt.ylabel('Total Spend (Scaled)')
plt.legend(title='Churn', loc='upper left', labels=['No Churn', 'Churn'])

# Scatter Plot 3: Usage Frequency vs. Support Calls
plt.subplot(1, 3, 3)
sns.scatterplot(x='Usage Frequency', y='Support Calls', hue='Churn', data=train_df, palette='viridis', alpha=0.6)
plt.title('Usage Frequency vs. Support Calls by Churn Status')
plt.xlabel('Usage Frequency (Scaled)')
plt.ylabel('Support Calls (Scaled)')
plt.legend(title='Churn', loc='upper left', labels=['No Churn', 'Churn'])

plt.tight_layout()
plt.show()
```



Interpretation of Scatter Plots:

The scatter plots provide a more detailed view of the relationships between specific pairs of features, colored by churn status.

- **Tenure vs. Total Spend by Churn Status**: This plot clearly shows a region where most churned customers (orange/red dots) are located in the upper-right quadrant (higher Tenure, higher Total Spend), while non-churned customers (green dots) are more evenly distributed across the lower-left quadrant.
- **Payment Delay vs. Total Spend by Churn Status**: Churned customers tend to have higher `Payment Delay` values, which is reflected in the vertical distribution of points in the plot.
- **Usage Frequency vs. Support Calls by Churn Status**: This plot is interesting because both features are positively correlated. Churned customers tend to have higher usage frequency and more support calls than non-churned customers.

These scatter plots highlight specific customer segments that are more prone to churn and provide visual evidence for our findings.

```
from sklearn.linear_model import LinearRegression

# Select 'Tenure' as the key numerical predictor
X_simple_lr_train = X_train[['Tenure']]
X_simple_lr_val = X_val[['Tenure']]

# Instantiate LinearRegression model
linear_model = LinearRegression()

# Train the model
linear_model.fit(X_simple_lr_train, y_train)

print("Linear Regression model trained successfully with 'Tenure' as the predictor.")
print(f"Model coefficients: {linear_model.coef_[0]:.4f}")
print(f"Model intercept: {linear_model.intercept_:.4f}")

Linear Regression model trained successfully with 'Tenure' as the predictor.
Model coefficients: -0.0257
Model intercept: 0.5668
```

```
X_test_lr = test_df[['Tenure']]

# Generate predictions on the validation set
y_val_pred_lr = linear_model.predict(X_simple_lr_val)

# Generate predictions on the test set
y_test_pred_lr = linear_model.predict(X_test_lr)

print("Predictions generated for validation and test sets.")

Predictions generated for validation and test sets.
```

```
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# Extract 'Churn' target from test_df for evaluation
y_test = test_df['Churn']

# Evaluate on Validation Set
r2_val = r2_score(y_val, y_val_pred_lr)
mae_val = mean_absolute_error(y_val, y_val_pred_lr)
mse_val = mean_squared_error(y_val, y_val_pred_lr)

print("Validation Set Evaluation:")
print(f" R-squared: {r2_val:.4f}")
print(f" Mean Absolute Error (MAE): {mae_val:.4f}")
print(f" Mean Squared Error (MSE): {mse_val:.4f}")

# Evaluate on Test Set
r2_test = r2_score(y_test, y_test_pred_lr)
mae_test = mean_absolute_error(y_test, y_test_pred_lr)
mse_test = mean_squared_error(y_test, y_test_pred_lr)

print("\nTest Set Evaluation:")
print(f" R-squared: {r2_test:.4f}")
print(f" Mean Absolute Error (MAE): {mae_test:.4f}")
print(f" Mean Squared Error (MSE): {mse_test:.4f}")
```

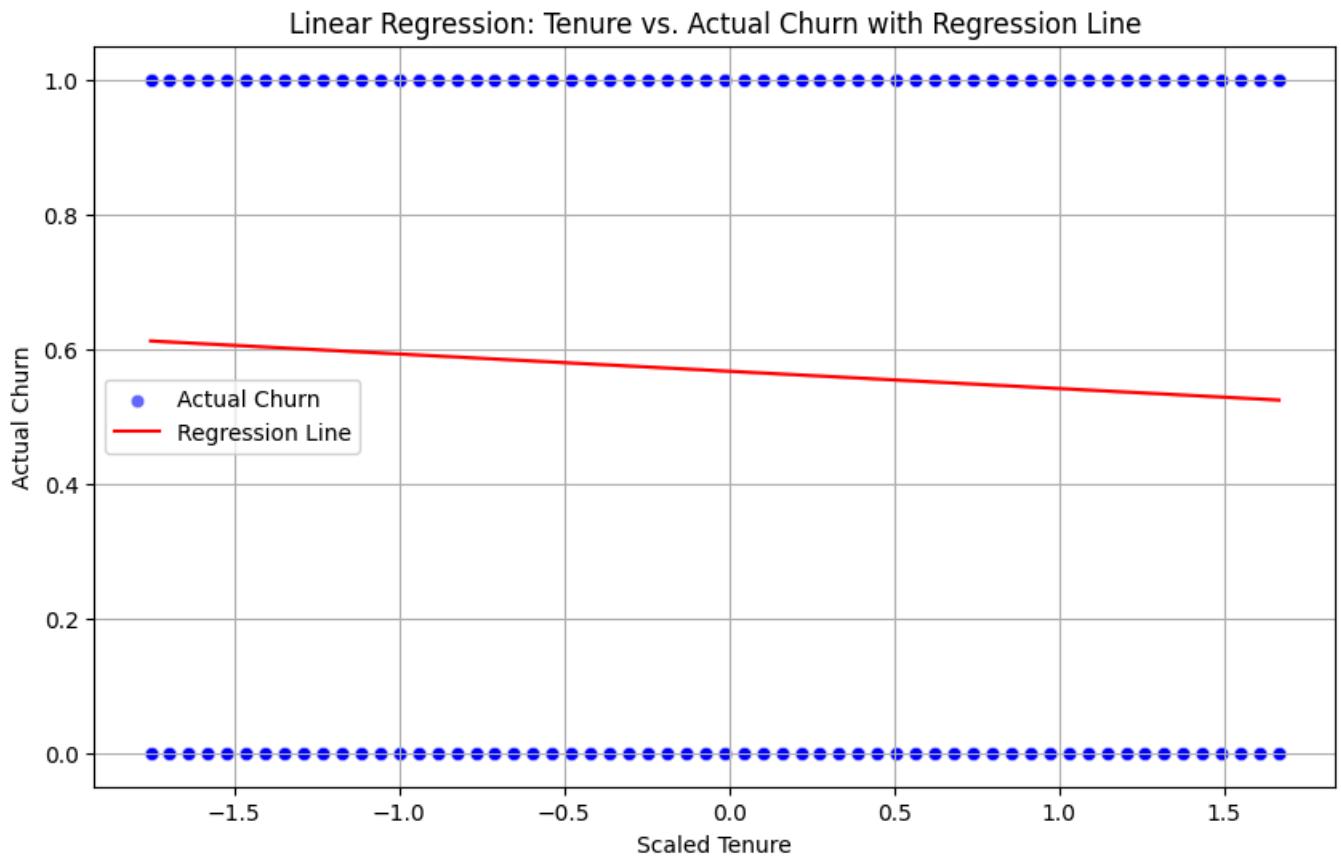
Validation Set Evaluation:
R-squared: 0.0028
Mean Absolute Error (MAE): 0.4895
Mean Squared Error (MSE): 0.2447

Test Set Evaluation:
R-squared: -0.0565
Mean Absolute Error (MAE): 0.5084
Mean Squared Error (MSE): 0.2634

```

plt.figure(figsize=(10, 6))
sns.scatterplot(x=X_simple_lr_val['Tenure'], y=y_val, label='Actual Churn', color='blue', alpha=0.6)
sns.lineplot(x=X_simple_lr_val['Tenure'], y=y_val_pred_lr, color='red', label='Regression Line')
plt.title('Linear Regression: Tenure vs. Actual Churn with Regression Line')
plt.xlabel('Scaled Tenure')
plt.ylabel('Actual Churn')
plt.legend()
plt.grid(True)
plt.show()

```



```

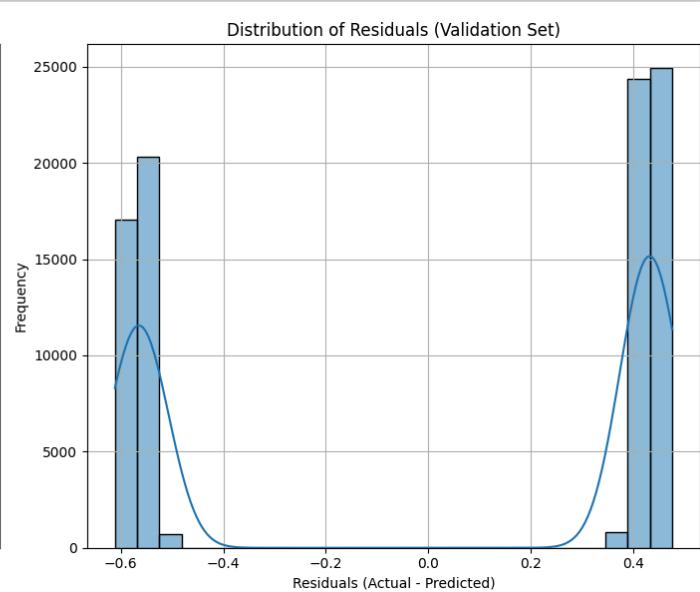
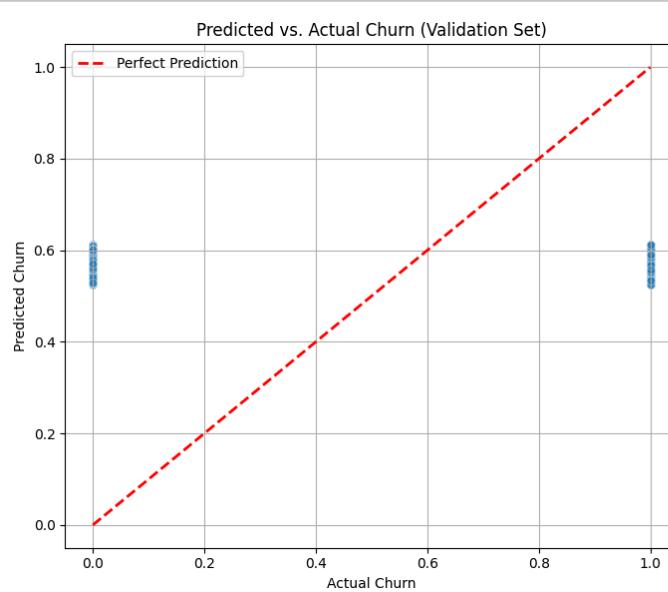
plt.figure(figsize=(14, 6))

# Predicted vs Actual Churn values
plt.subplot(1, 2, 1)
sns.scatterplot(x=y_val, y=y_val_pred_lr, alpha=0.6)
plt.plot([y_val.min(), y_val.max()], [y_val.min(), y_val.max()], 'r--', lw=2, label='Perfect Prediction')
plt.title('Predicted vs. Actual Churn (Validation Set)')
plt.xlabel('Actual Churn')
plt.ylabel('Predicted Churn')
plt.legend()
plt.grid(True)

# Distribution of Residuals
residuals = y_val - y_val_pred_lr
plt.subplot(1, 2, 2)
sns.histplot(residuals, kde=True)
plt.title('Distribution of Residuals (Validation Set)')
plt.xlabel('Residuals (Actual - Predicted)')
plt.ylabel('Frequency')
plt.grid(True)

plt.tight_layout()
plt.show()

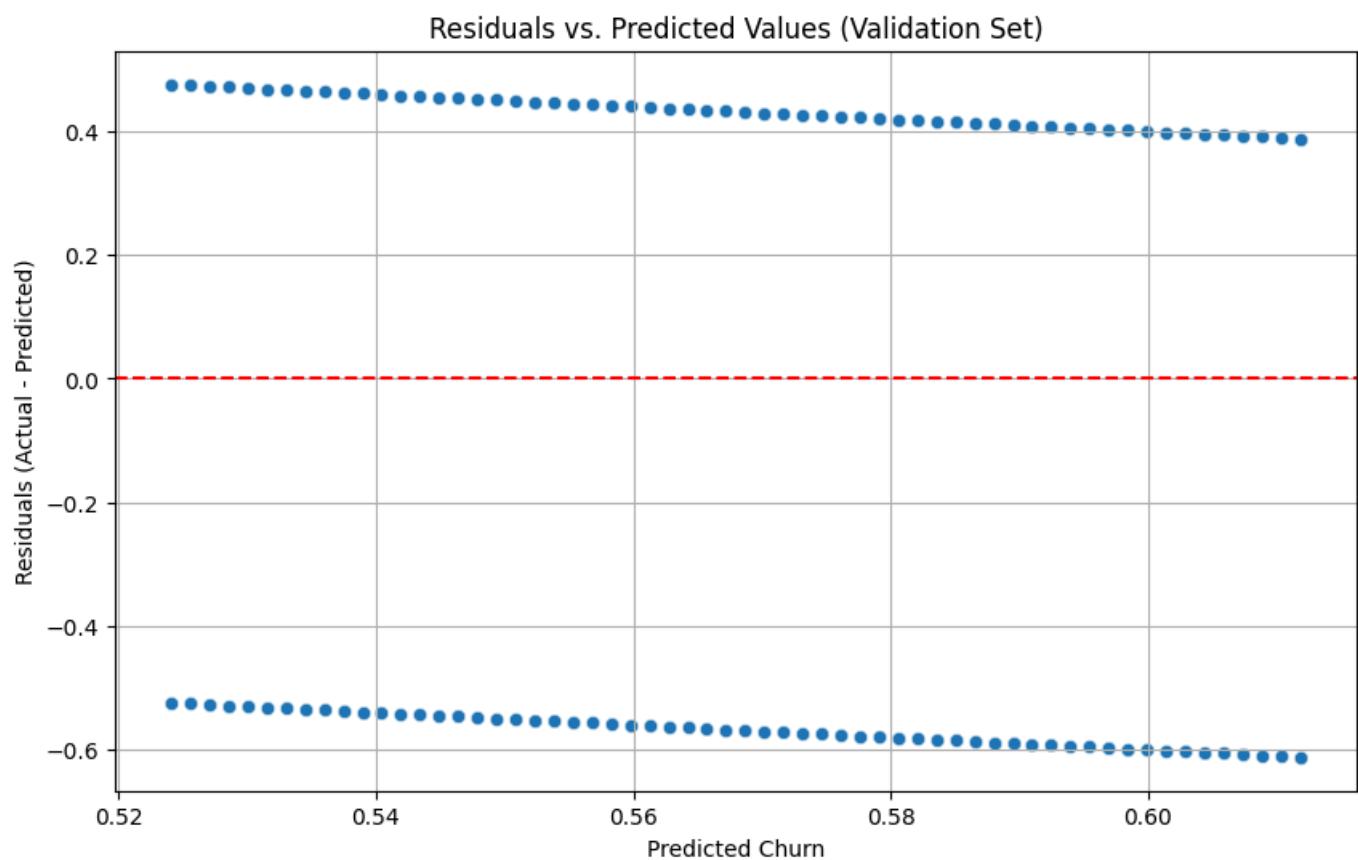
```



```

plt.figure(figsize=(10, 6))
sns.scatterplot(x=y_val_pred_lr, y=residuals, alpha=0.6)
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residuals vs. Predicted Values (Validation Set)')
plt.xlabel('Predicted Churn')
plt.ylabel('Residuals (Actual - Predicted)')
plt.grid(True)
plt.show()

```



```

def reconstruct_categorical(df, prefix):
    cols = [col for col in df.columns if col.startswith(prefix)]

```

```

if cols:
    # Find the column that is True (1) for each row and extract its suffix
    return df[cols].idxmax(axis=1).apply(lambda x: x.split('_')[1])
return None

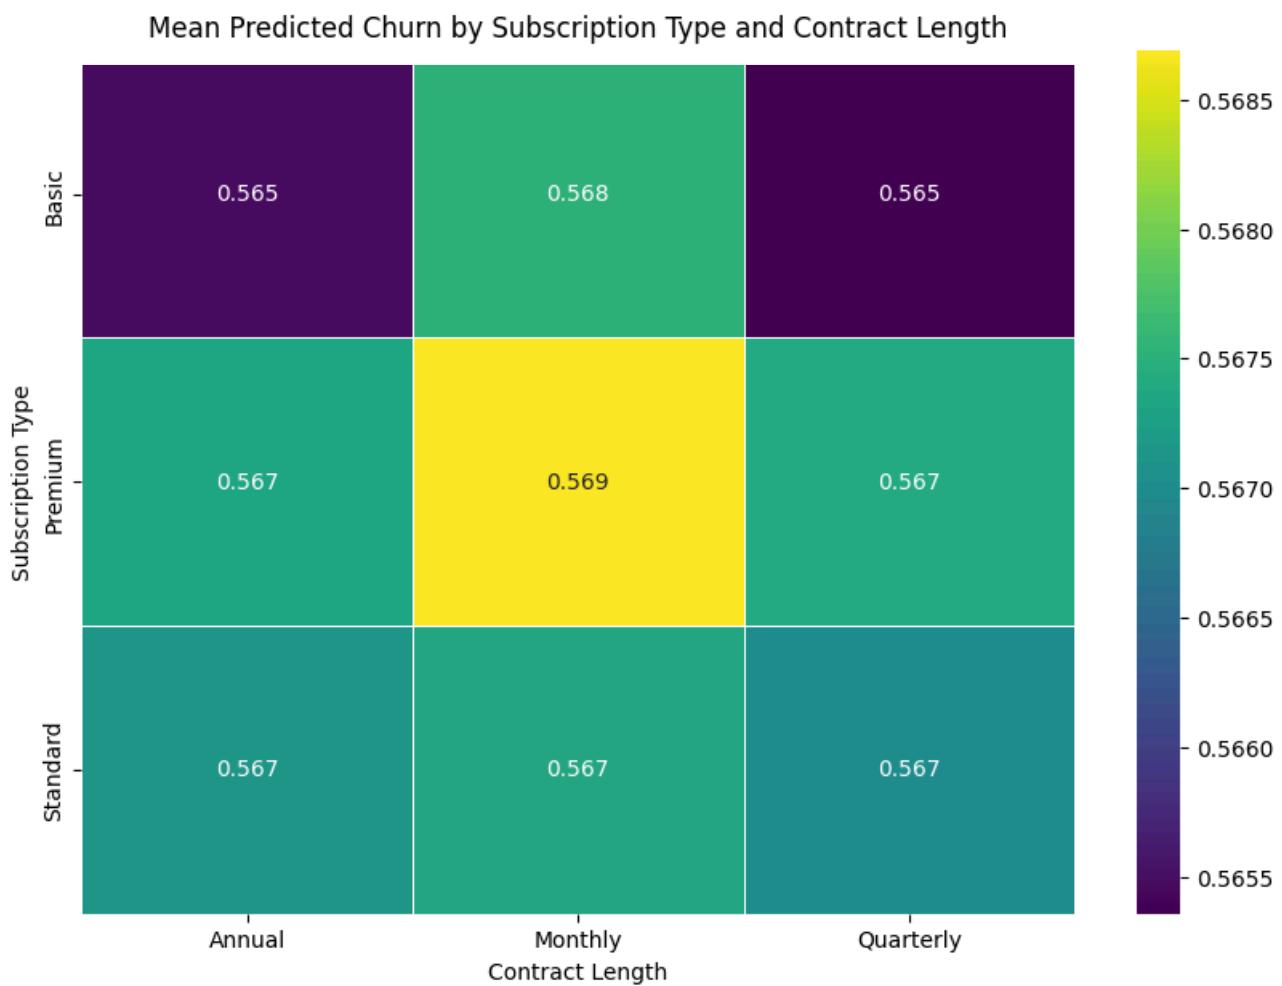
# Create a temporary DataFrame for validation set analysis
val_analysis_df = X_val.copy()
val_analysis_df['Predicted_Churn'] = y_val_pred_lr

# Reconstruct original categorical columns for grouping
val_analysis_df['Subscription Type'] = reconstruct_categorical(val_analysis_df, 'Subscription Type_')
val_analysis_df['Contract Length'] = reconstruct_categorical(val_analysis_df, 'Contract Length_')

# Create a pivot table to get the mean predicted churn for each group
heatmap_data = val_analysis_df.pivot_table(index='Subscription Type', columns='Contract Length', values='Predicted_Churn')

plt.figure(figsize=(10, 7))
sns.heatmap(heatmap_data, annot=True, cmap='viridis', fmt=".3f", linewidths=.5)
plt.title('Mean Predicted Churn by Subscription Type and Contract Length')
plt.ylabel('Subscription Type')
plt.xlabel('Contract Length')
plt.show()

```



Interpretation of Simple Linear Regression Model:

The simple linear regression model was trained using 'Tenure' as the sole predictor for 'Churn'.

- Model Coefficients:** The coefficient for 'Tenure' was approximately **-0.0257**, and the intercept was around **0.5668**.
- Trend Interpretation:**
 - The negative coefficient for 'Tenure' indicates an inverse relationship between 'Tenure' and 'Churn'. Specifically, for every one-unit increase in **Scaled Tenure**, the predicted 'Churn' probability **decreases by approximately 0.0257**.
 - The intercept of 0.5668 can be interpreted as the predicted churn probability when **Scaled Tenure** is zero. Since 'Churn' is a binary variable (0 or 1), this can be seen as a baseline probability.

- This trend aligns with our EDA findings, where longer tenure was associated with a lower likelihood of churn. However, given that churn is a binary outcome, a linear regression attempts to fit a straight line through points that are only 0 or 1, which isn't ideal.
- **Implications for Predicting Churn:** The model suggests that `Tenure` is a statistically significant predictor, as its coefficient is non-zero. The negative sign confirms that customers with longer relationships (higher tenure) are less likely to churn. However, the magnitude of the coefficient (-0.0257) is relatively small, suggesting that `Tenure` alone might not be a very strong predictor of individual churn probability in a linear sense. The model will predict continuous values for churn, which then need to be thresholded to get a binary classification. This can lead to issues, as the predicted values can fall outside the meaningful range of 0 to 1.

Limitations of Simple Linear Regression for Binary Classification:

Using a simple linear regression model for a binary classification target like 'Churn' (0 for no churn, 1 for churn) comes with several inherent limitations, which are clearly reflected in the evaluation metrics:

- **Non-binary Predictions:** Linear regression models predict continuous numerical values. When applied to a binary target, the model can output predictions outside the meaningful range of [0, 1]. For example, the model might predict a 'churn probability' of -0.2 or 1.3, which are nonsensical in the context of probability. To convert these to a binary outcome, a threshold must be applied, which is an additional step not inherently handled by the model.
- **Assumptions Violation:** Linear regression assumes a linear relationship between the independent and dependent variables, and that the residuals are normally distributed with constant variance (homoscedasticity). For a binary dependent variable, these assumptions are almost always violated. The relationship is inherently non-linear (an S-curve for probabilities), and the residuals will not be normally distributed or homoscedastic (they will concentrate at 0 and 1).
- **Poor Performance Metrics:**
 - **Low R-squared (0.0028 for validation, -0.0565 for test):** R-squared measures the proportion of variance in the dependent variable that is predictable from the independent variable(s). A very low or negative R-squared value, as observed here, indicates that 'Tenure' explains almost none of the variability in 'Churn'. A negative R-squared suggests that the model performs worse than simply predicting the mean of the dependent variable.
 - **High Mean Absolute Error (MAE) and Mean Squared Error (MSE):** While MAE (0.4895 validation, 0.5084 test) and MSE (0.2447 validation, 0.2634 test) provide a sense of prediction error, their interpretation is less straightforward for a binary target. An MAE of ~0.5 indicates that, on average, the model's predictions are about 0.5 units away from the true binary labels. Given the target is 0 or 1, this means predictions are often quite far off.
- **Difficulty in Interpretation:** While the coefficient indicates a negative trend, interpreting the output as a 'probability' is inaccurate due to the continuous nature of the predictions and the violated assumptions. The model struggles to differentiate between the two classes effectively, often predicting values around the mean churn rate, as observed in the scatter plot of actual vs. predicted churn values.